
COMP9318: Data Warehousing and Data Mining

Assignment Project Exam Help

— L2: Data Warehousing and OLAP —
<https://powcoder.com>

Add WeChat powcoder

- Why and What are Data Warehouses?

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Data Analysis Problems

- The same data found in many different systems
 - Example: customer data across different departments
 - The same concept is defined differently
- Heterogeneous sources
 - Relational DBMS, OnLine Transaction Processing (OLTP)
 - Unstructured data in files (e.g., MS Excel) and documents (e.g., MS Word)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Data Analysis Problems (Cont'd)

- Data is suited for operational systems
 - Accounting, billing, etc.
 - Do not support analysis across business functions
- Data quality is bad
 - Missing data, imprecise data, different use of systems
- Data are “volatile”
 - Data deleted in operational systems (6months)
 - Data change over time – no historical information

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Solution: Data Warehouse

- Defined in many different ways, but not rigorously.
 - A decision support database that is maintained **separately** from the organization's operational database
 - Support **information processing** by providing a solid platform of consolidated, historical data for analysis.
- "A data warehouse is a subject-oriented, time-variant, and nonvolatile collection of data in support of management's decision-making process."—W. H. Inmon
- Data warehousing:
 - The process of constructing and using data warehouses

Data Warehouse—Subject-Oriented

- Organized around major subjects, such as **customer, product, sales**.
- Focusing on the modeling and analysis of data for decision makers, not on **daily operations or transaction processing**.
- Provide **a simple and concise view around particular subject issues by excluding data that are not useful in the decision support process**.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Data Warehouse—Integrated

- Constructed by integrating multiple, heterogeneous data sources
 - relational databases, flat files, on-line transaction records
- Data cleaning and data integration techniques are applied.
 - Ensure consistency in naming conventions, encoding structures, attribute measures, etc. among different data sources
 - E.g., Hotel price: currency, tax, breakfast covered, etc.
 - When data is moved to the warehouse, it is converted.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Data Warehouse—Time Variant

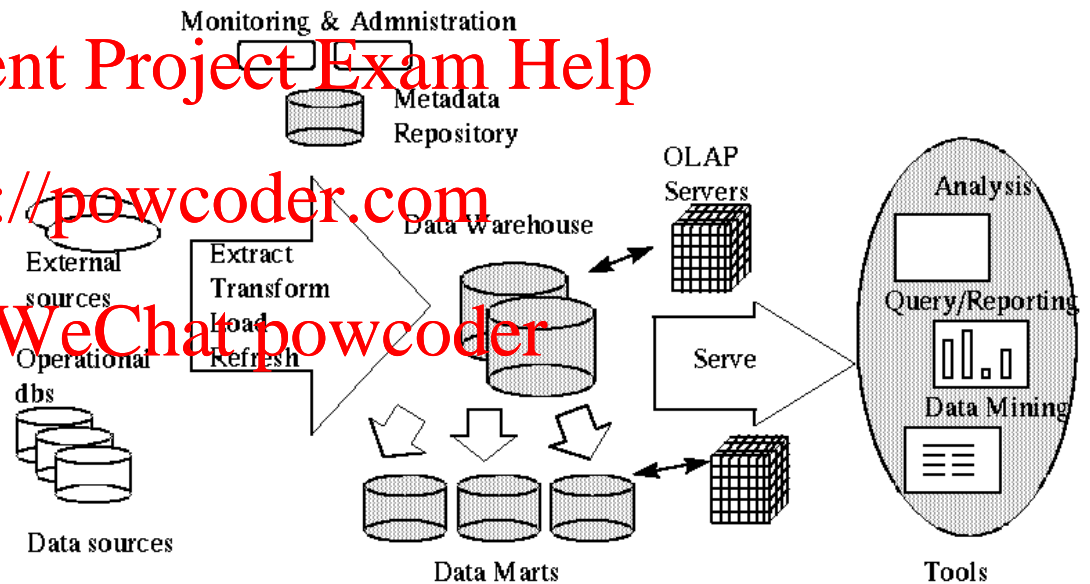
- The time horizon for the data warehouse is significantly longer than that of operational systems.
 - Operational database: current value data.
 - Data warehouse data: provide information from a historical perspective (e.g., past 5-10 years)
- Every key structure in the data warehouse
 - Contains an element of time, explicitly or implicitly
 - But the key of operational data may or may not contain “time element”.

Data Warehouse—Non-Volatile

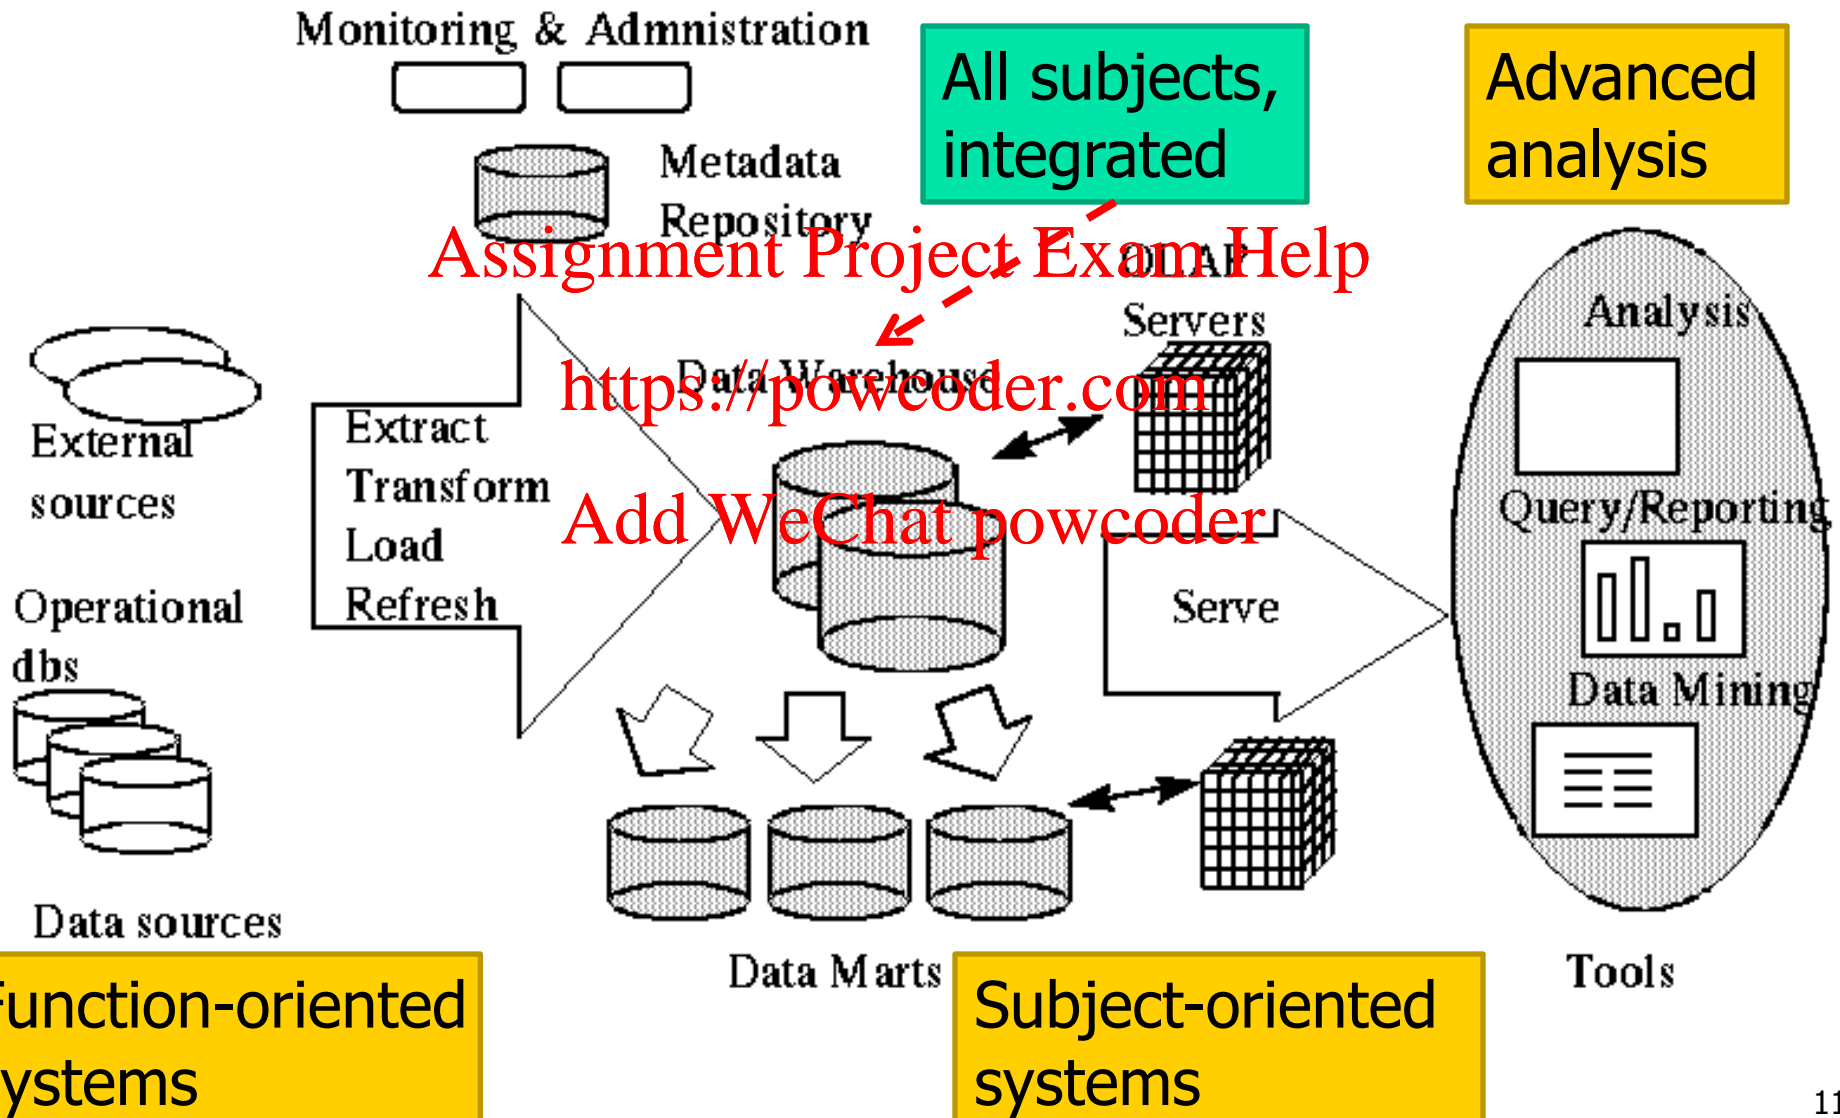
1. A **physically separate store** of data transformed from the operational environment.
2. Operational **updates of data do not help** in the data warehouse environment.
<https://powcoder.com>
 - Does not require transaction processing, recovery, and concurrency control mechanisms
 - Requires only two operations in data accessing:
 - *initial loading of data* and *access of data*.

Data Warehouse Architecture

- Extract data from operational data sources
 - clean, transform
- Bulk load/refresh
 - warehouse is offline
- OLAP-server provides multidimensional view
- Multidimensional-olap (Essbase, oracle express)
- Relational-olap (Redbrick, Informix, Sybase, SQL server)



Data Warehouse Architecture



Why Separate Data Warehouse?

- High performance for both systems
 - DBMS—tuned for OLTP: access methods, indexing, concurrency control, recovery
 - Warehouse—tuned for OLAP: complex OLAP queries, multidimensional view, consolidation
- Different functions and different data:
 - missing data: Decision support requires historical data which operational DBs do not typically maintain
 - data consolidation: DS requires consolidation (aggregation, summarization) of data from heterogeneous sources
 - data quality: different sources typically use inconsistent data representations, codes and formats which have to be reconciled

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Why OLAP Servers?

- Different workload:
 - OLTP (on-line transaction processing)
 - Major task of traditional relational DBMS
 - Day-to-day operations: purchasing, inventory, banking, manufacturing, payroll, registration, accounting, etc.
 - OLAP (on-line analytical processing)
 - Major task of data warehouse system
 - Data analysis and decision making
- Queries hard/infeasible for OLTP, e.g.,
 - Which **week** we have the largest sales?
 - Does the sales of **dairy products** increase over time?
 - Generate a **spread sheet** of total sales by state and by year.
- Difficult to represent these queries by using SQL ← Why?

OLTP vs. OLAP

	OLTP	OLAP
users	clerk, IT professional	knowledge worker
function	day to day operations	decision support
DB design	application-oriented	subject-oriented
data	current, up-to-date detailed, flat relational isolated	historical, summarized, multidimensional integrated, consolidated
usage	repetitive	ad-hoc
access	read/write index/hash on prim. key	lots of scans
unit of work	short, simple transaction	complex query
# records accessed	tens	millions
#users	thousands	hundreds
DB size	100MB-GB	100GB-TB
metric	transaction throughput	query throughput, response

Comparisons

	Databases	Data Warehouses
Purpose	Many purposes; Flexible and general	One purpose: Data analysis
Conceptual Model	ER	Multidimensional
Logical Model	(Normalized) Relational Model	(Denormalized) Star schema / Data cube/cuboids
Physical Model	Relational Tables	ROLAP: Relational tables MOLAP: Multidimensional arrays
Query Language	SQL (hard for analytical queries)	MDX (easier for analytical queries)
Query Processing	B+-tree/hash indexes, Multiple join optimization, Materialized views	Bitmap/Join indexes, Star join, Materialized data cube

- The Multidimensional Model

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

The Multidimensional Model

- A data warehouse is based on a multidimensional data model which views data in the form of a **data cube**, which is a multidimensional generalization of 2D spread sheet.
- Key concepts:
 - **Facts**: the subject it models
 - Typically transactions in this course; other types includes snapshots, etc.
 - Measures: numbers that can be aggregated
 - Dimensions: context of the measure
 - Hierarchies:
 - Provide contexts of different granularities (aka. grains)
- Goals for dimensional modeling:
 - Surround facts with as much relevant context (dimensions) as possible ← Why?

Supermarket Example

- Subject: analyze total sales and profits
- Fact: Each Sales **Transaction**
 - Measure: Dollars_Sold, Amount_Sold, Cost
 - Calculated Measure: Profit
- Dimensions: <https://powcoder.com>
 - Store [Add WeChat powcoder](#)
 - Product
 - Time

Visualizing the Cubes

- A valid **instance** of the model is a data cube

total Sales		product			
city		p1	p2	p3	p4
	NY	\$454	-	-	\$925
	LA	\$468	\$800	-	-
	SD	\$296	-	\$240	-
	SF	\$652	-	\$540	\$745

		product			
city		p1	p2	p3	p4
	454				925
	468	800			
	296		240		
	652		540	745	

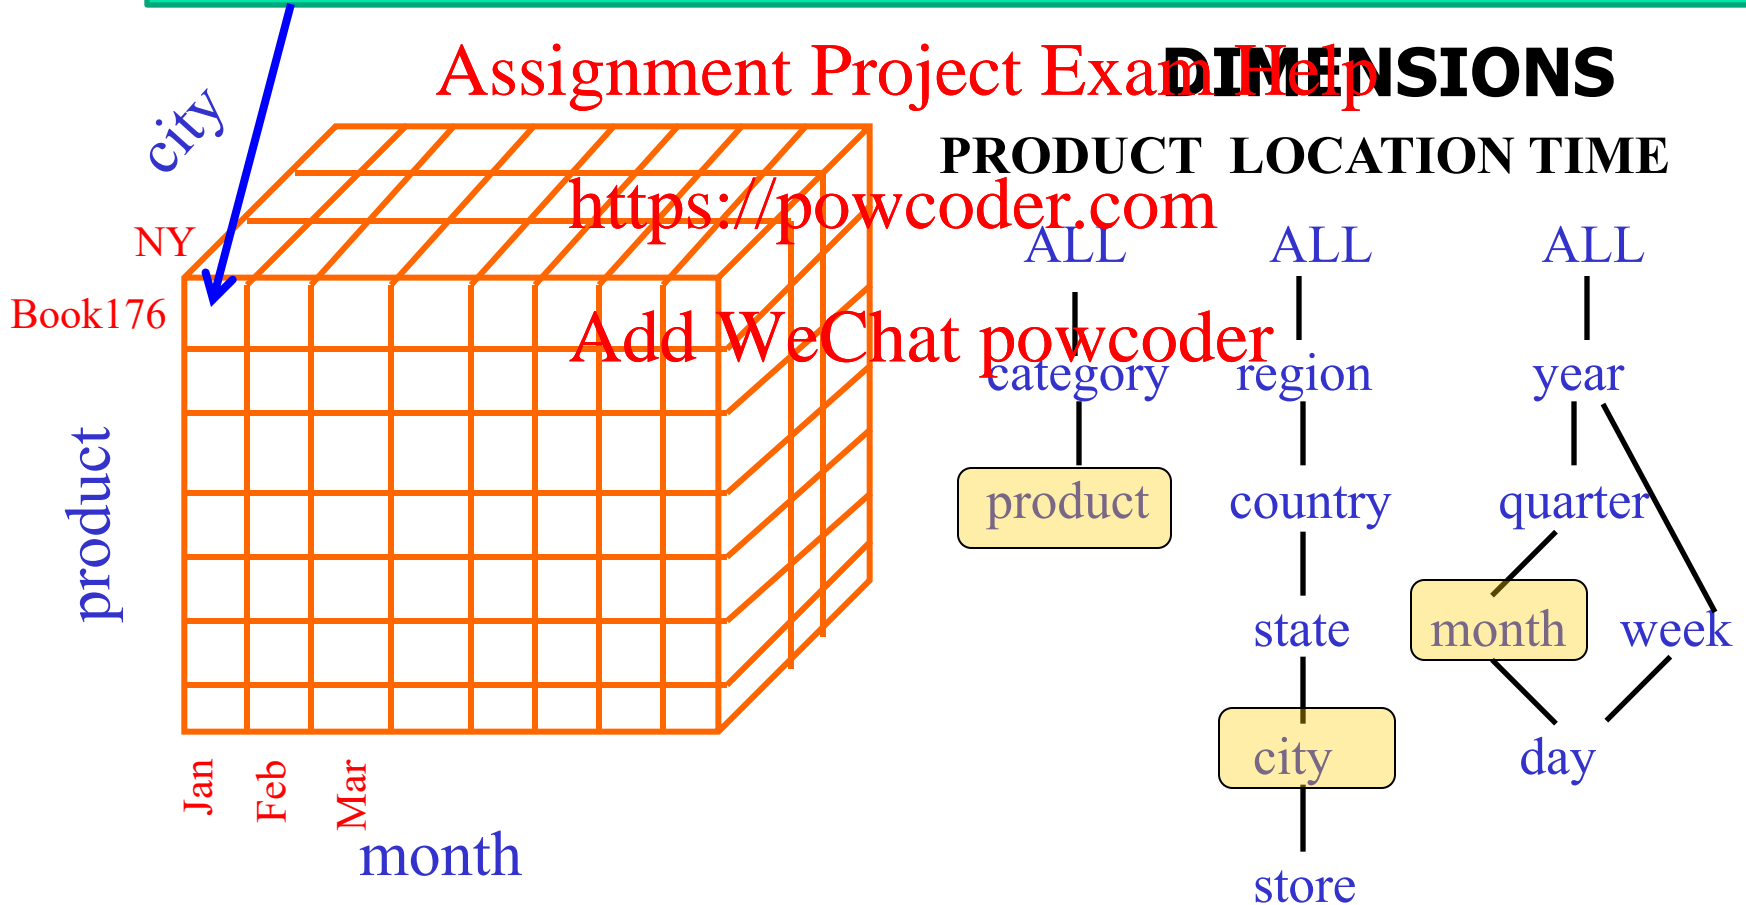
Concepts: cell, fact (=non-empty cell), measure, dimensions

Q: How to generalize it to 3D?

3D Cube and Hierarchies

Concepts: hierarchy (a tree of dimension values), level

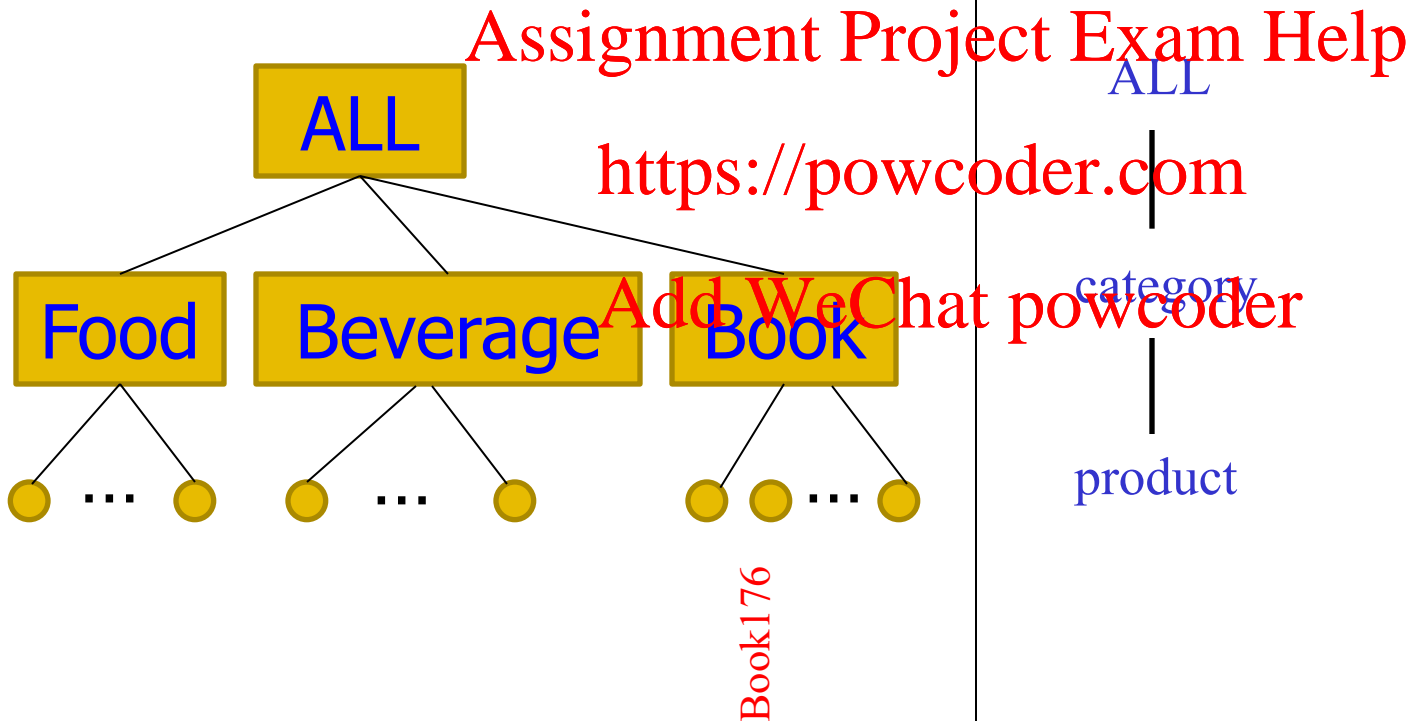
Sales of **book176** in **NY** in **Jan** can be found in this cell



Hierarchies


Concepts: hierarchy (a tree of dimension values), level

Which design is better? Why?



The (city, moth) Cuboid

Sales of **ALL_PROD** in **NY** in **Jan**



DIMENSIONS
PRODUCT LOCATION TIME

ALL

category

product

ALL

region

country

state

city

store

ALL

year

quarter

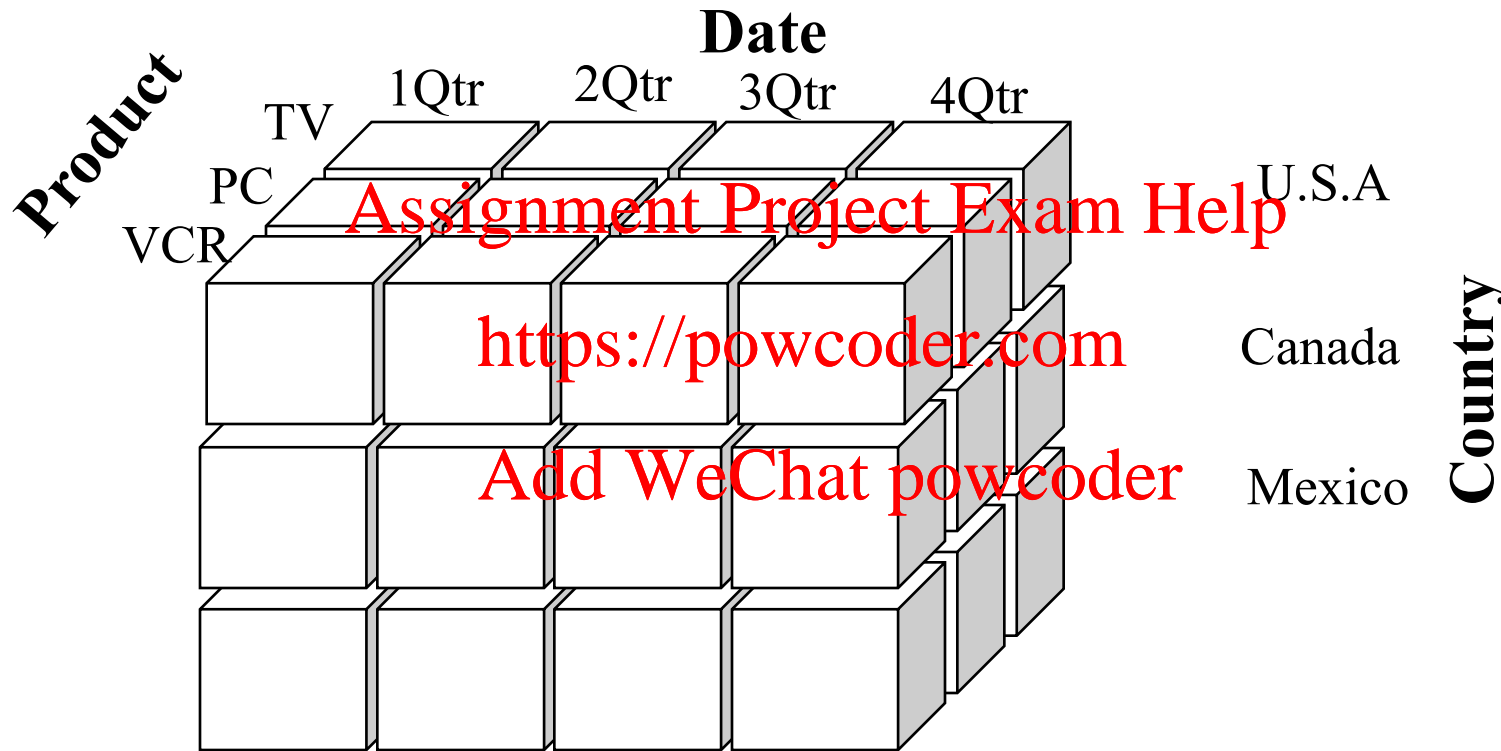
month

day

week


All the Cuboids

Assume: no other non-ALL levels on all dimensions.

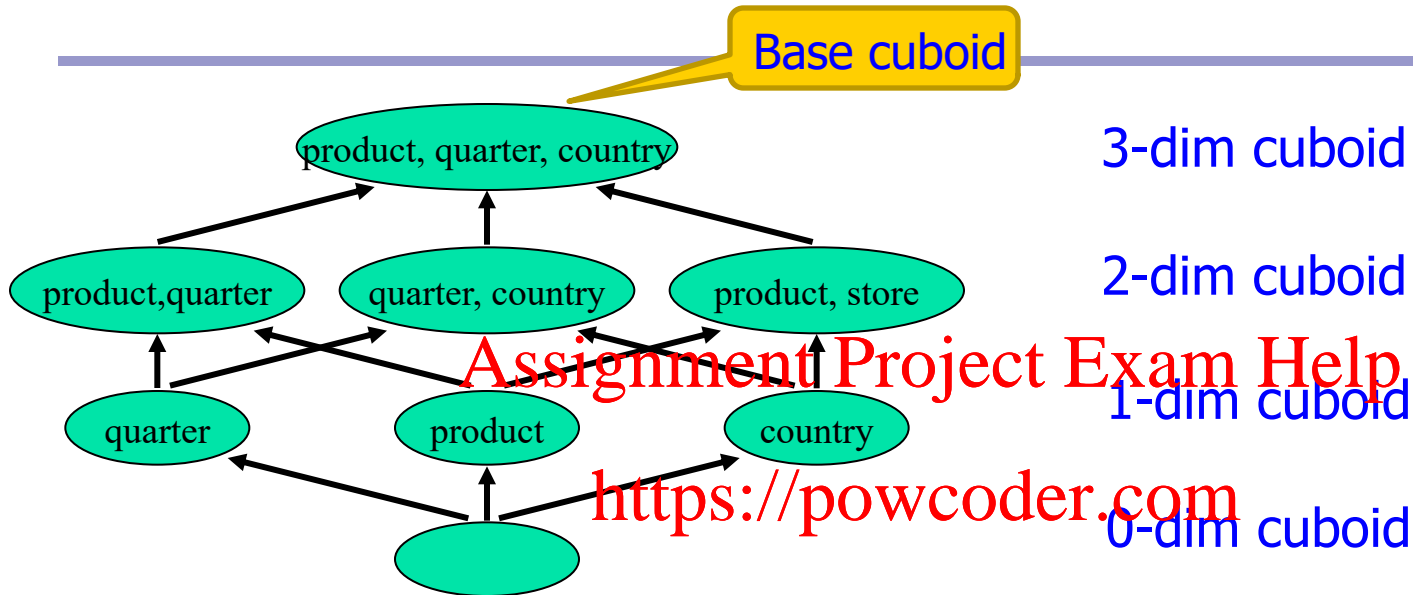


All the Cuboids /2

Assume: no other non-ALL levels on all dimensions.



Lattice of the cuboids



- n-dim cube can be represented as (D_1, D_2, \dots, D_n) , where D_i is the set of allowed values on the i-th dimension.
 - if $D_i = L_i$ (a particular level), then $D_i =$ all descendant dimension values of L_i .
 - ALL can be omitted and hence reduces the effective dimensionality
- A complete cube of d-dimensions consists of $\prod_{i=1}^d (n_i + 1)$ cuboids, where n_i is the number of levels (excluding ALL) on i-th dimension.
 - They collectively form a lattice.

Properties of Operations

- All operations are closed under the multidimensional model
 - i.e., both input and output of an operation is a cube
 - So that they can be composed
- <https://powcoder.com>
Add WeChat powcoder

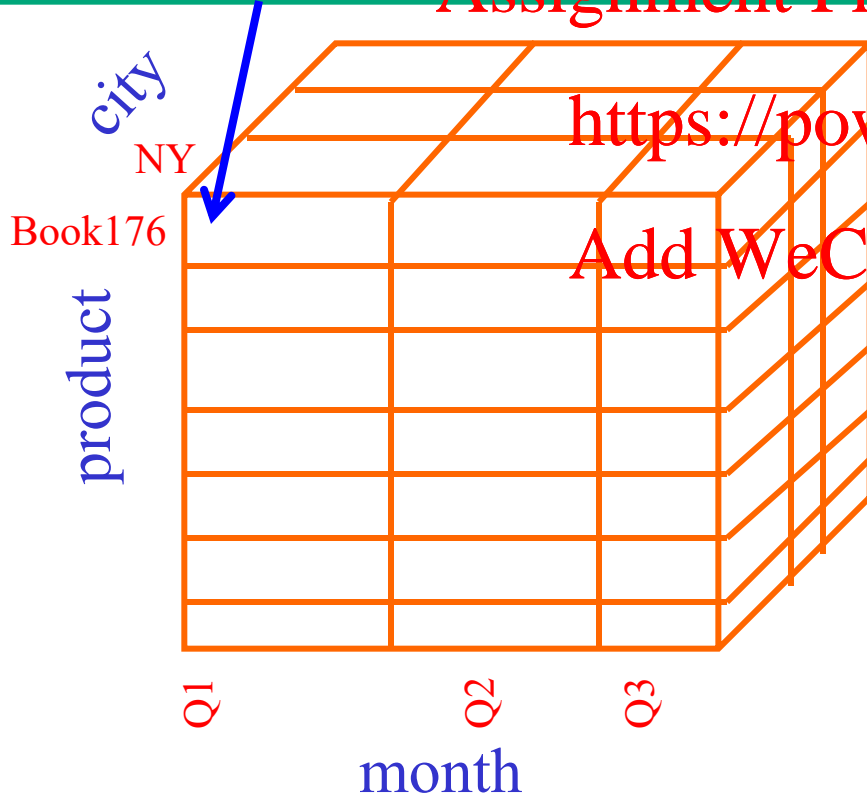
Q: What's the analogy in the Relational Model?

Common OLAP Operations

- **Roll-up**: move up the hierarchy

Q: what should be its value?

Sales of **book176** in **NY** in **Q1** here



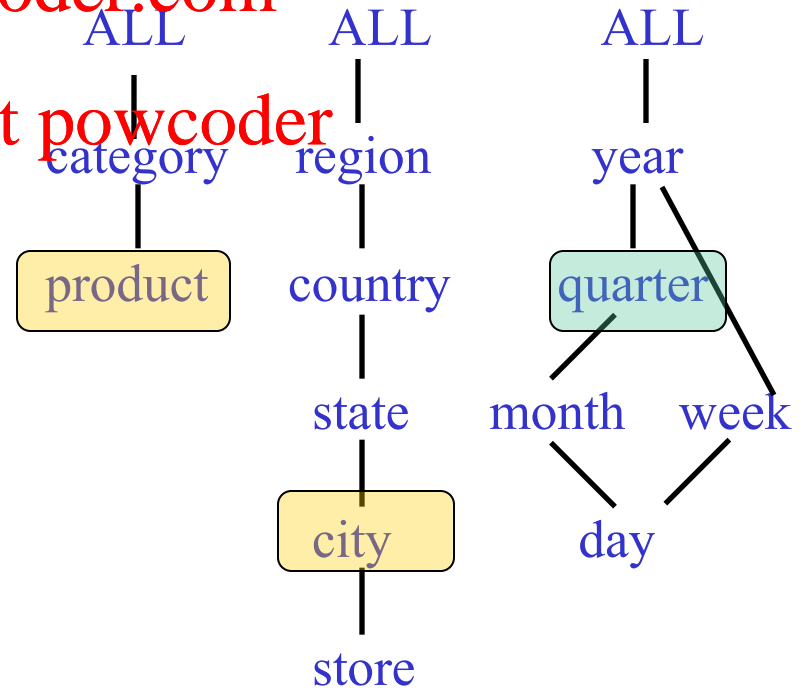
Assignment Project Exam Help



<https://powcoder.com>

Add WeChat powcoder

DIMENSIONS

PRODUCT LOCATION TIME

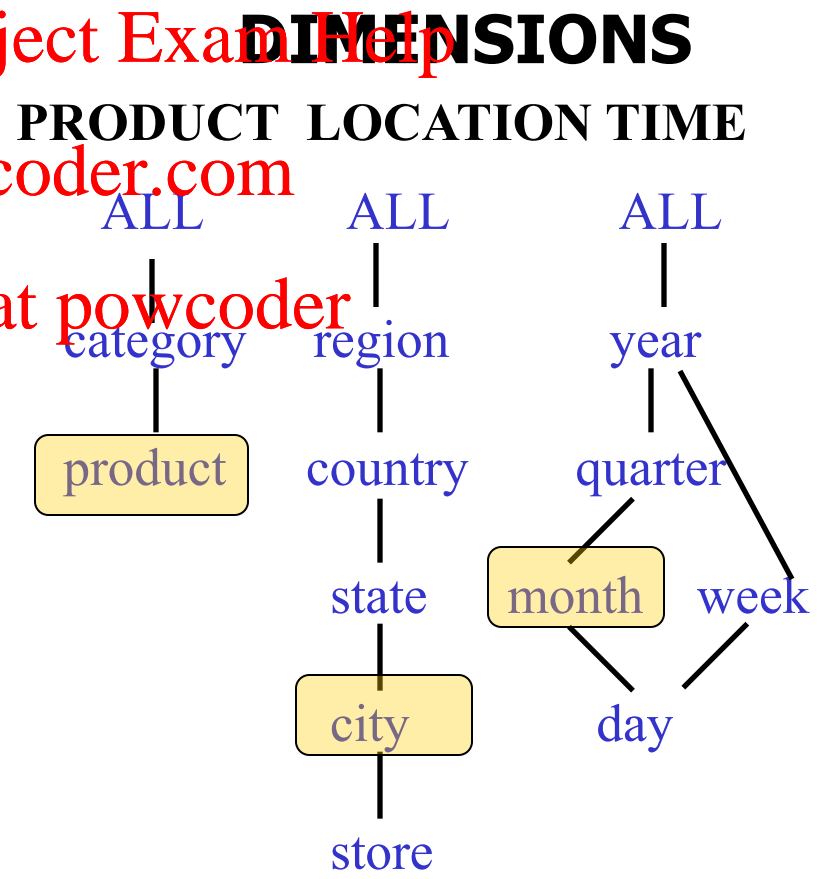




Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Data Cube Measures: Three Categories

- Distributive: if the result derived by applying the function to n aggregate values is the same as that derived by applying the function on all the data without partitioning
 - E.g., count(), sum(), min(), max()
- Algebraic: if it can be computed by an algebraic function with M arguments (where M is a bounded integer), each of which is obtained by applying a distributive aggregate function
 - E.g., avg(), min_N(), standard_deviation()
- Holistic: if there is no constant bound on the storage size needed to describe a subaggregate.
 - E.g., median(), mode(), rank()

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Common OLAP Operations

- **Drill-down**: move down the hierarchy

- more fine-grained aggregation

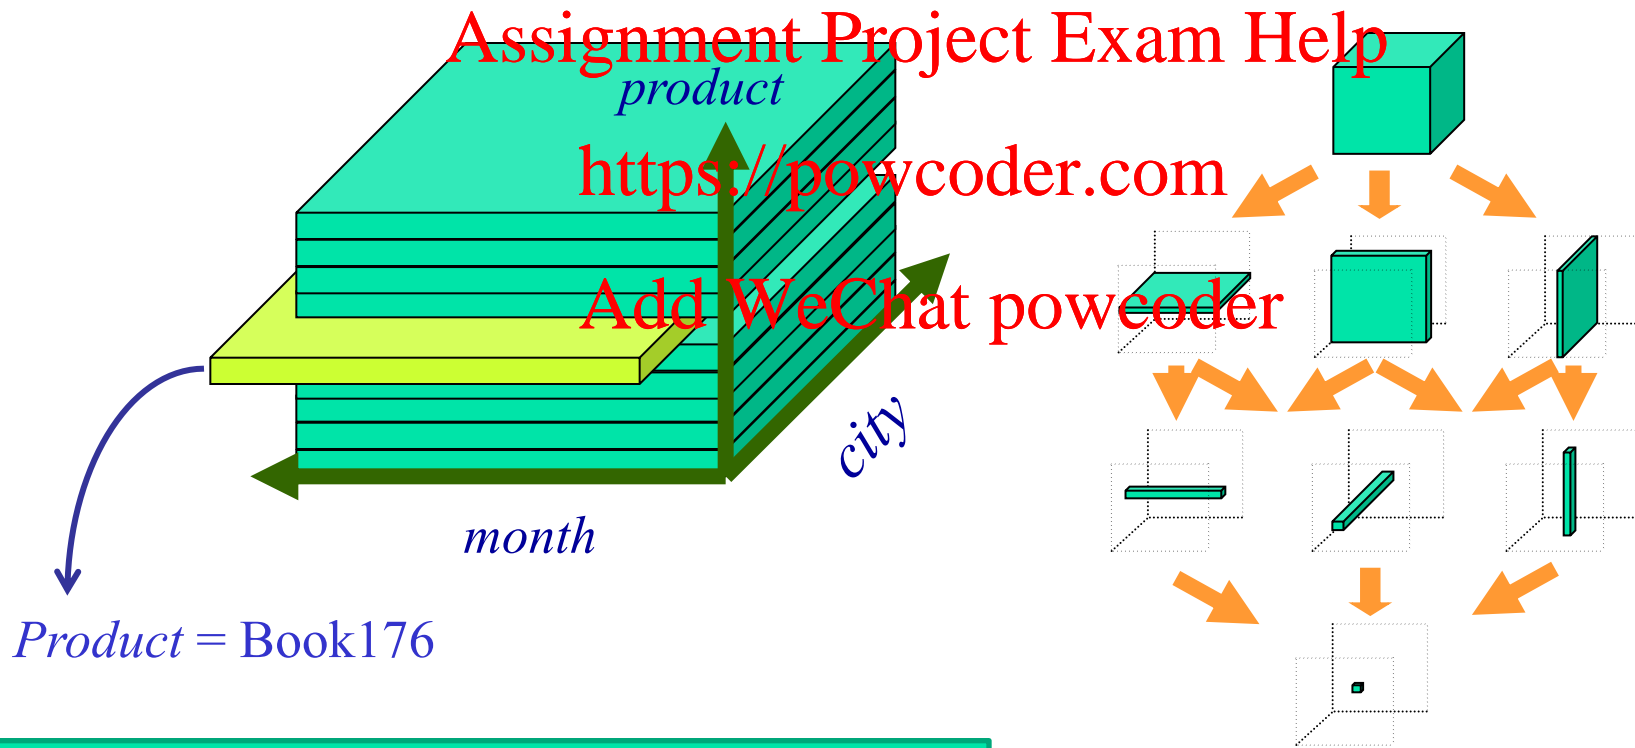
Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Slice and Dice Queries


- Slice and Dice: select and project on one or more dimension values



The output cube has smaller dimensionality than the input cube

Pivoting

- Pivoting: aggregate on selected dimensions
 - usually 2 dims (cross-tabulation)




Assignment Project Exam Help

Sales (of all products) in
NY in Q1
=sum(???)

<https://powcoder.com>

Add WeChat powcoder

pivot on (city, month)



City
	

NY
	Q1	Q2	Q3
	month		

A Reflective Pause

- Let's review the definition of data cubes again.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- **Key message:**
 - Disentangle the “object” from its “representation” or “implementation”

Modeling Exercise 1: Monthly Phone Service Billing

The screenshot shows the Telstra website homepage. At the top is a blue header with the Telstra logo on the left, a search bar on the right, and a navigation menu in the center containing links for 'Telstra Home', 'Shop Online', 'My Account', 'Products & Services', 'Help Centre', and 'About Telstra'. Below the header, the page is divided into several sections. On the left, there is a 'My Account' section with a 'Login to My Account' button, and a 'Quick Links' section with a list of links: 'Check my WebMail', 'Pay my bill', 'Find a mobile', 'Move home', 'Activate Pre-Paid', 'Search FAQs', 'Switch to Telstra', 'Find a Telstra or T[life] shop', 'Discover Next G™', and 'Contact Us'. Below this is a 'Latest News' section. The main content area features a large blue banner for 'Whereis® Mobile' with a mobile phone image and text about free browsing on Next G™ mobile. To the right of this is another banner for 'For Next G™ handset coverage questions' with a 'Find out more' button. Below the banners, there are two columns of service tiles. The 'Personal' column includes 'Mobile', 'BigPond® Internet', 'Bundle & Save' (showing a phone, a laptop, and a TV), 'Home Phone', and 'FOXTEL AUSTAR'. The 'Business' column includes 'Small & Medium Business' and 'Enterprise & Government'. A 'Login to Online Services' button is also present. A large red watermark is overlaid across the center of the page, reading 'Assignment Project Exam Help' and 'https://powcoder.com'. Another red watermark, 'Add WeChat powcoder', is also visible.

Theme: analyze the income/revenue of Telstra

Solution

- FACT

Assignment Project Exam Help

- MEASURE <https://powcoder.com>

Add WeChat powcoder

- DIMENSIONS

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- The Logical Model

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Logical Models

- Two main approaches:
 - Using relational DB technology:
 - Star schema, Snowflake schema, Fact constellation
 - Using multidimensional technology:
 - Just as multidimensional data cube

<https://powcoder.com>

Add WeChat powcoder

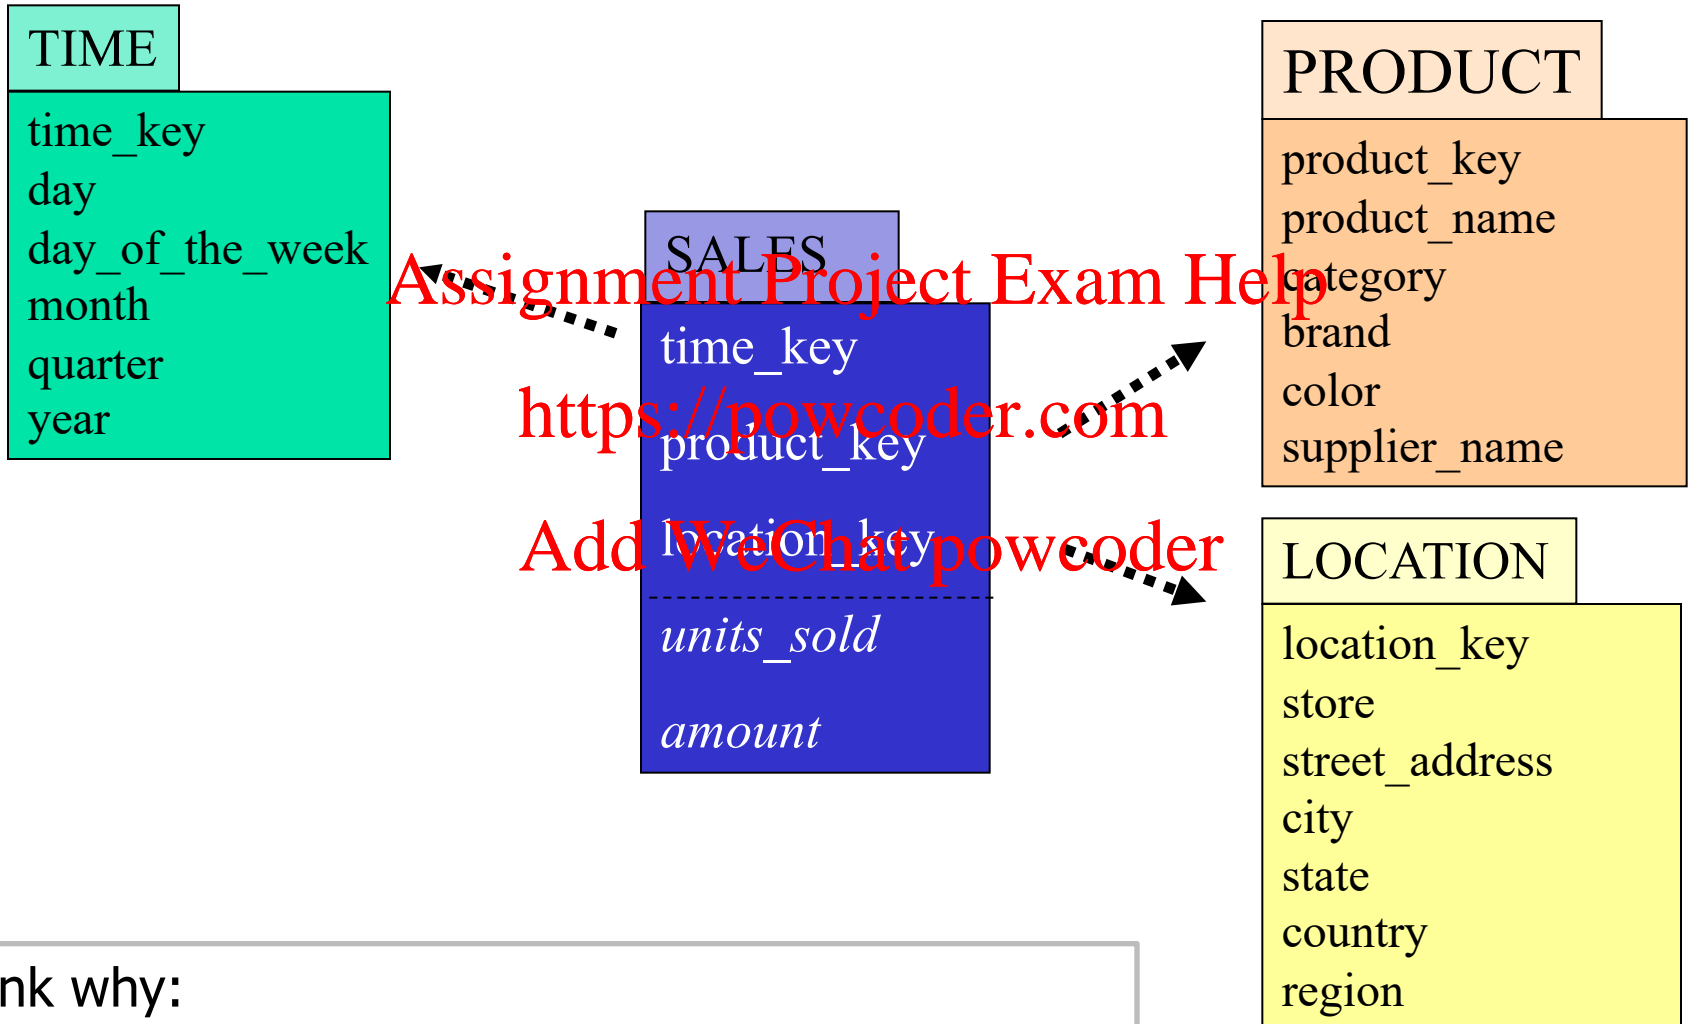
Universal Schema → Star Schema

- Many data warehouses adopt a star schema to represent the multidimensional model
- Each dimension is represented by a dimension-table
 - LOCATION (location_key, Store, street, address, city, state, country, region)
 - dimension tables are not normalized
- Transactions are described through a fact-table
 - each tuple consists of a logical pointer to each of the dimension-tables (foreign-key) and a list of measures (e.g. sales \$\$\$)

The universal schema for supermarket

Store	City	State	Prod	Brand	Category	\$Sold	#Sold	Cost
S136	Syd	NSW	76Ha	Nestle	Biscuit	40	10	18
S173	Melb	Vic	76Ha	Nestle	Biscuit	20	5	11

The Star Schema



Think why:

- (1) Denormalized **once** from the universal schema
- (2) Controlled **redundancy**

Typical Models for Data Warehouses

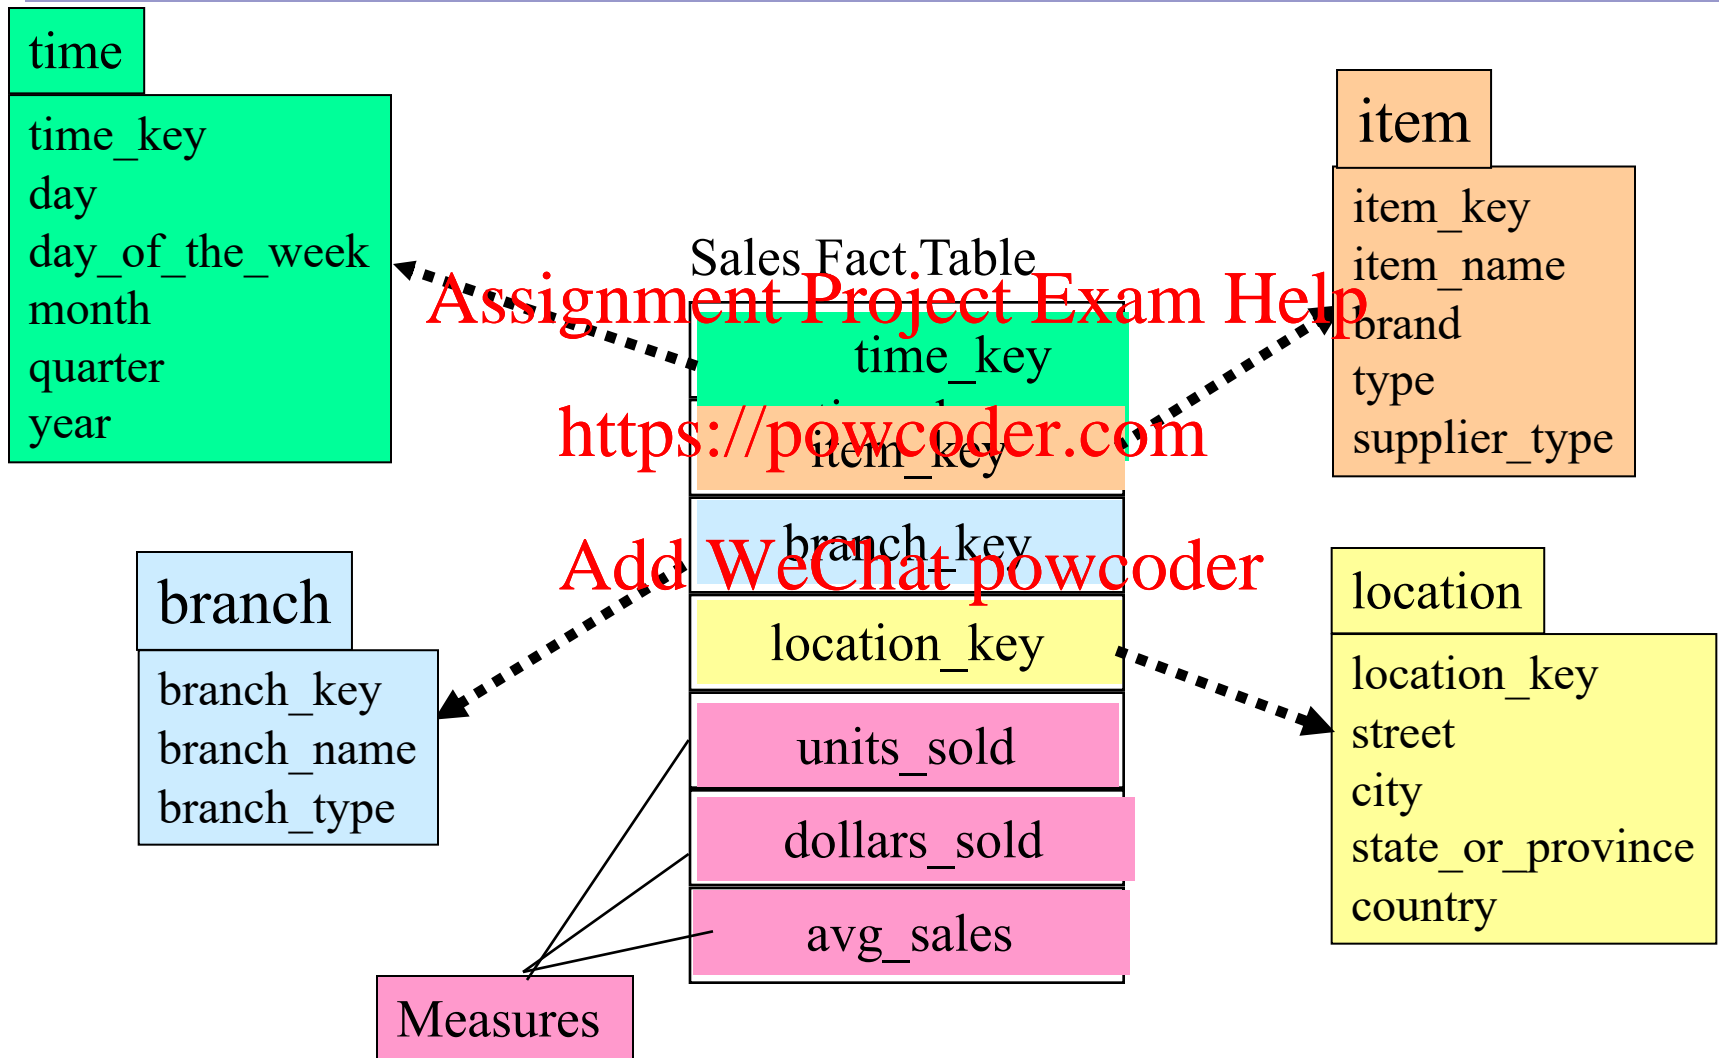
- Modeling data warehouses: dimensions & measures
 - Star schema: A fact table in the middle connected to a set of dimension tables
 - Snowflake schema: A refinement of star schema where some dimensional hierarchy is normalized into a set of smaller dimension tables, forming a shape similar to snowflake
 - Fact constellations: Multiple fact tables share dimension tables, viewed as a collection of stars, therefore called galaxy schema or fact constellation

Assignment Project Exam Help

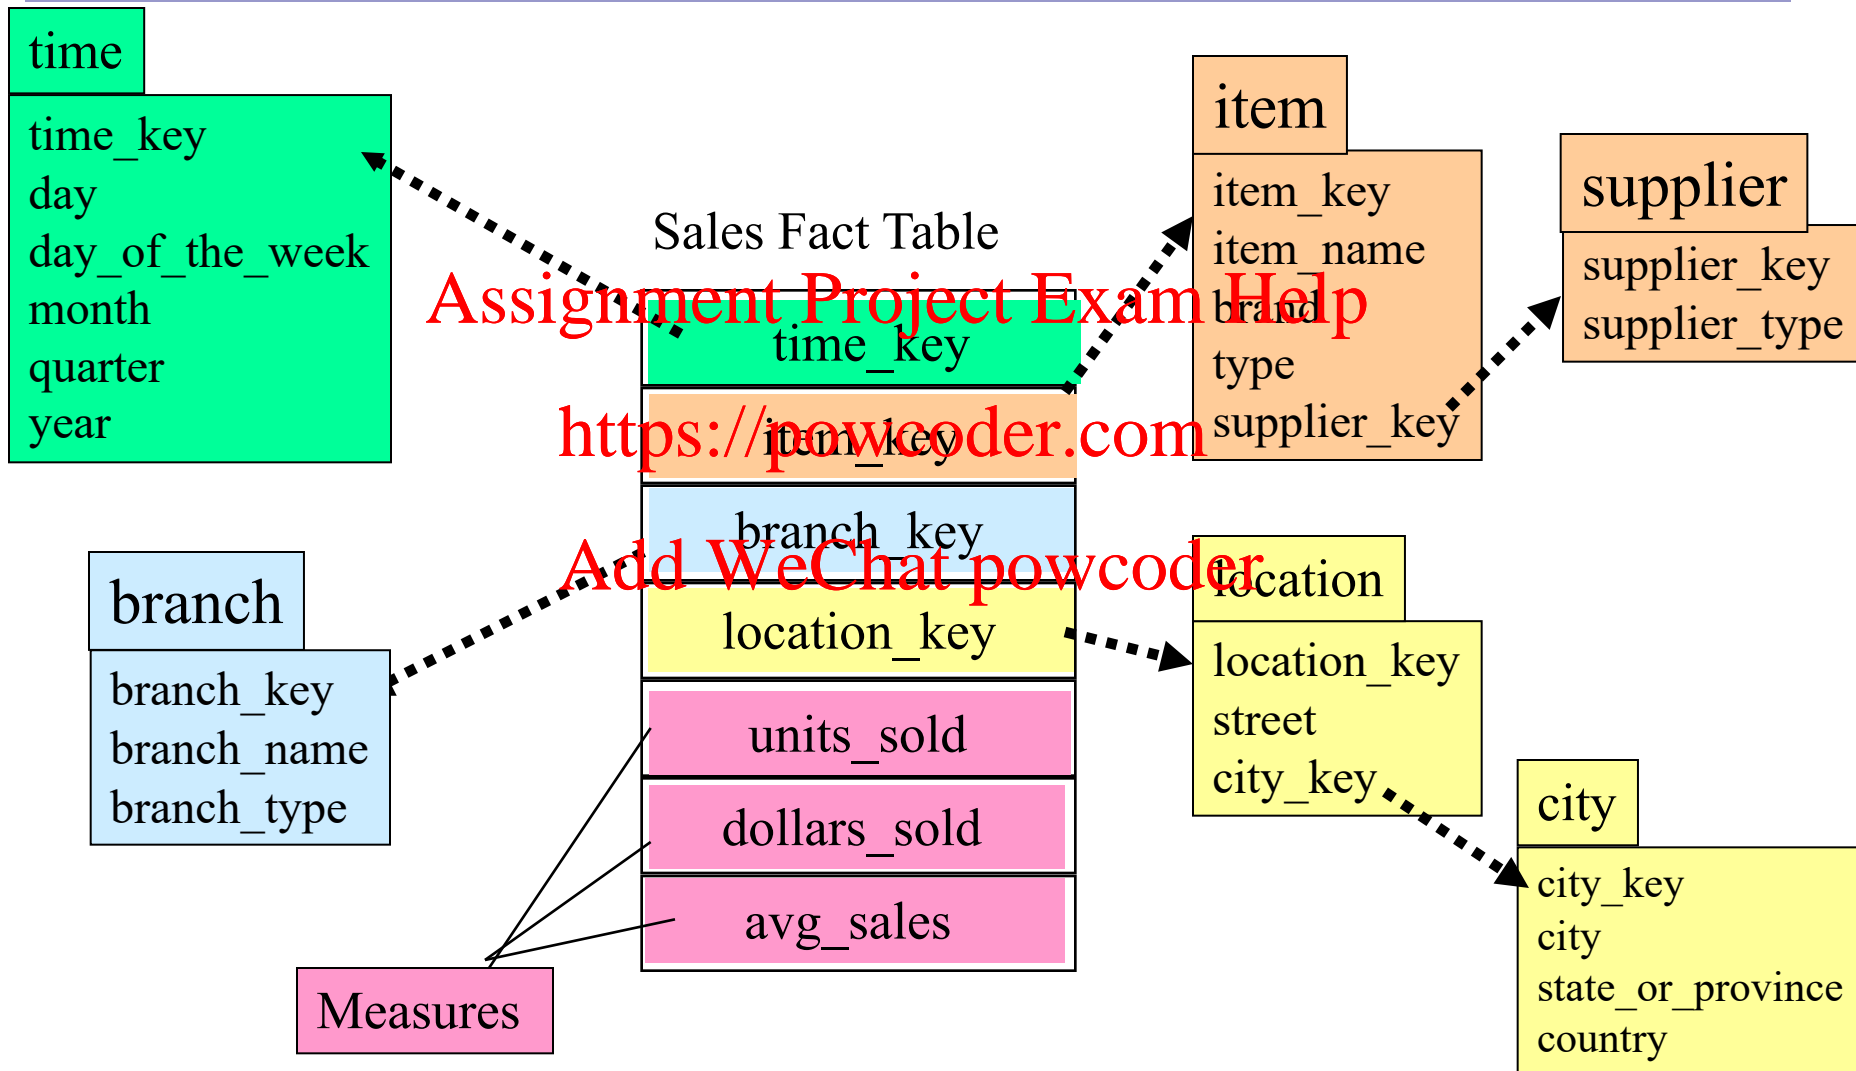
<https://powcoder.com>

Add WeChat powcoder

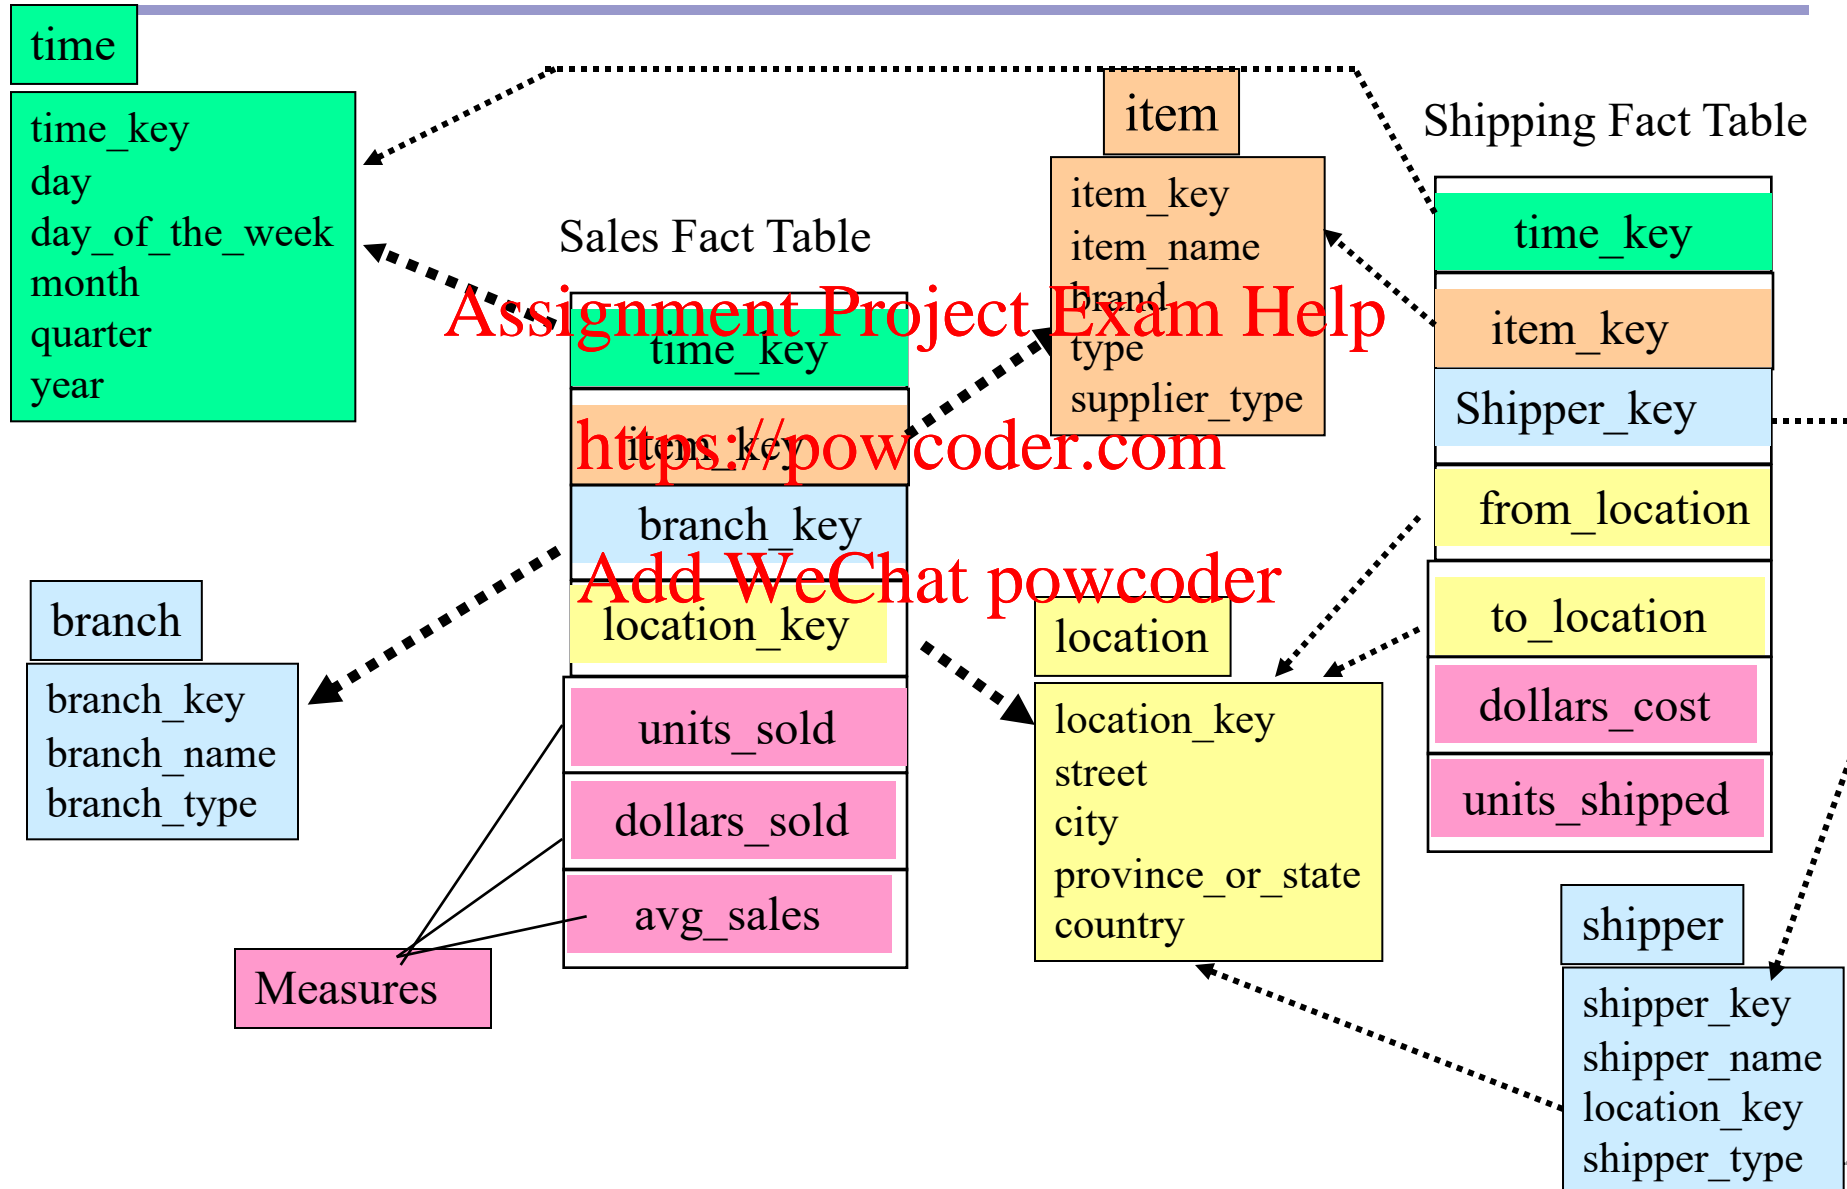
Example of Star Schema



Example of Snowflake Schema



Example of Fact Constellation



Advantages of Star Schema

- Facts and dimensions are clearly depicted
 - dimension tables are relatively static, data is loaded (append mostly) into fact table(s)
 - easy to comprehend (and write queries)

“Find total sales per product category in our stores in Europe”

```
SELECT PRODUCT.category, SUM(SALES.amount)
FROM    SALES, PRODUCT, LOCATION
WHERE    SALES.product_key = PRODUCT.product_key
AND      SALES.location_key = LOCATION.location_key
AND      LOCATION.region="Europe"
GROUP BY PRODUCT.category
```

Operations: Slice (Loc.Region.Europe) + Pivot (Prod.category)

- Query Language

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Query Language


- Two approaches:
 - Using relational DB technology: SQL (with extensions such as CUBE/PIVOT/UNPIVOT)
 - Using multidimensional technology: MDX

Assignment Project Exam Help

<https://powcoder.com>

```
SELECT PRODUCT.category,  
SUM(SALES.amount)  
FROM SALES, PRODUCT, LOCATION  
WHERE SALES.product_key =  
PRODUCT.product_key  
AND SALES.location_key =  
LOCATION.location_key  
AND LOCATION.region="Europe"  
GROUP BY PRODUCT.category
```

```
SELECT  
{[PRODUCT].[category]} on ROWS,  
{[MEASURES].[amount]} on COLUMNS  
FROM [SALES]  
WHERE ([LOCATION].[region].[Europe])
```



Add WeChat powcoder

Operations: Slice (Loc.Region.Europe) + Pivot (Prod.category, Measures.amnt)

- Physical Model + Query Processing Techniques

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Physical Model + Query Processing Techniques

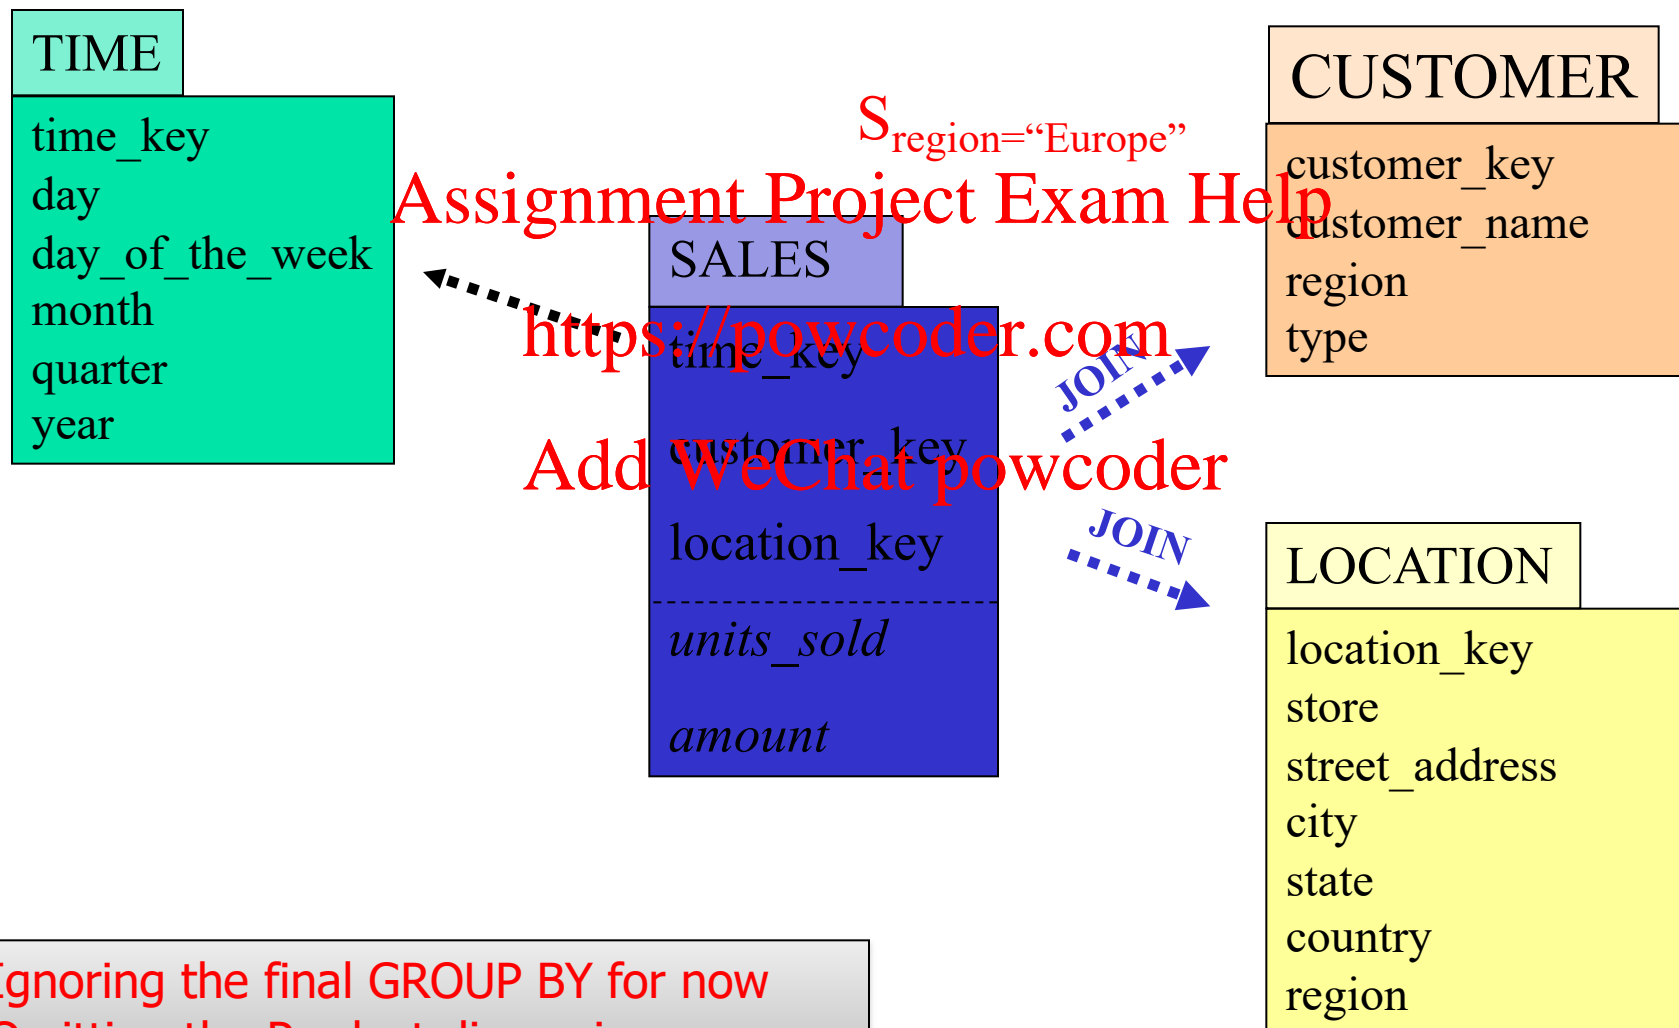
- Two main approaches:
 - Using relational DB technology: ROLAP
 - Using multidimensional technology: MOLAP
- Hybrid: HOLAP
 - Base cuboid: ROLAP
 - Other cuboids: MOLAP

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Q1: Selection on low-cardinality attributes



- Ignoring the final GROUP BY for now
- Omitting the Product dimension

Indexing OLAP Data: Bitmap Index

(1) BI on dimension tables

- Index on an attribute (column) with low distinct values
- Each distinct values, v , is associated with a n -bit vector ($n = \#rows$)
 - The i -th bit is set if the i -th row of the table has the value v for the indexed column
- Multiple BIs can be efficiently combined to enable optimized scan of the table

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Custom

Cust	Region	Type
C1	Asia	Retail
C2	Europe	Dealer
C3	Asia	Dealer
C4	America	Retail
C5	Europe	Dealer

BI on Customer.Region

v	bitmap
Asia	1 0 1 0 0
Europe	0 1 0 0 1
America	0 0 0 1 0

Indexing OLAP Data: Bitmap Index /2

- (1) Bitmap join index (BI on Fact Table Joined with Dimension tables)
- Conceptually, perform a join, map each dimension value to the bitmap of corresponding fact table rows.

<https://powcoder.com>

Add WeChat powcoder

-- ORACLE SYNTAX --

```
CREATE BITMAP INDEX sales_cust_region_bjix
ON   sales(customer.cust_region)
FROM sales, customer
WHERE sales.cust_id = customers.cust_id;
```

Indexing OLAP Data: Bitmap Index /3

Sales

time	customer	loc	Sale
101	C1	100	1
173	C1	200	2
208	C2	100	3
863	C3	200	5
991	C1	100	8
1001	C2	200	13
1966	C4	100	21
2017	C5	200	34

Customer

Cust	Region	Type
C1	Asia	Retail
C2	Europe	Dealer
C3	Asia	Dealer
C4	America	Retail
C5	Europe	Dealer

Assignment Project Exam Help

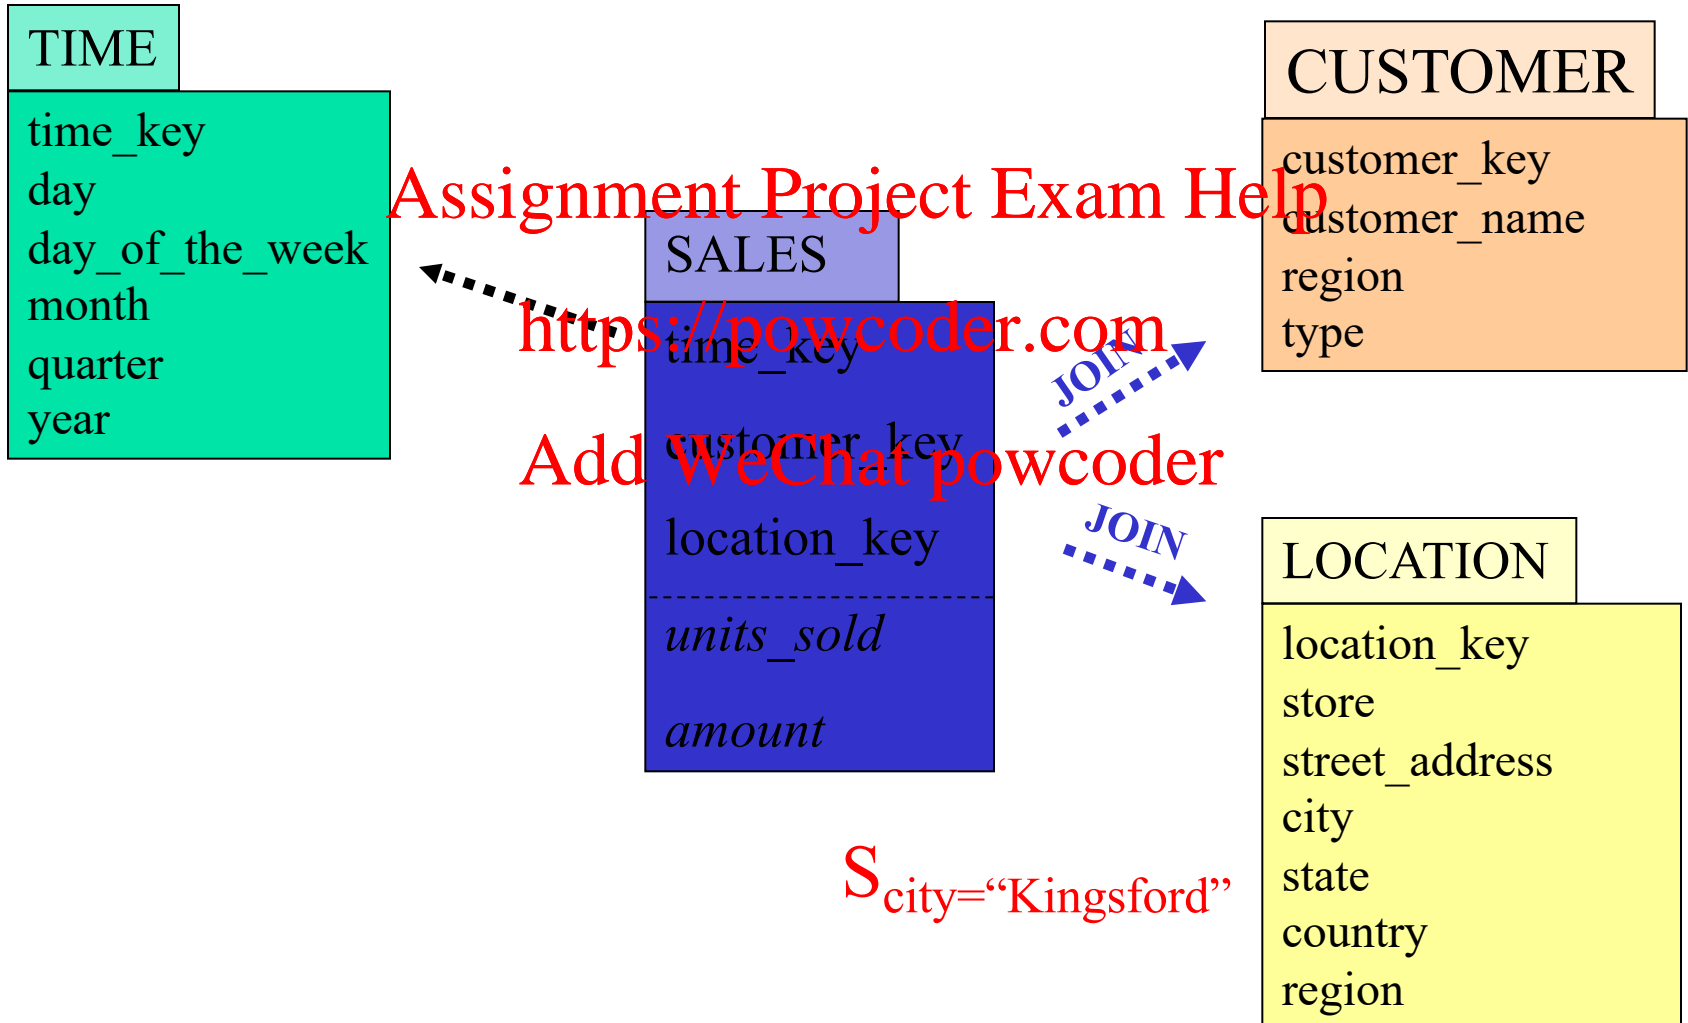
<https://powcoder.com>

Add WeChat powcoder

BI on Sales(Customer.Region)

v	bitmap
Asia	11011000
Europe	00100101
America	00000010

Q2: Selection on high-cardinality attributes



Indexing OLAP Data: Join Indices

- Join index relates the values of the dimensions of a star schema to rows in the fact table.
 - a join index on city maintains for each distinct city a list of ROW-IDs of the tuples recording the sales in the city
- Join indices can span multiple dimensions OR
 - can be implemented as bitmap-indexes (per dimension)
 - use bit-op for multiple-joins

Location

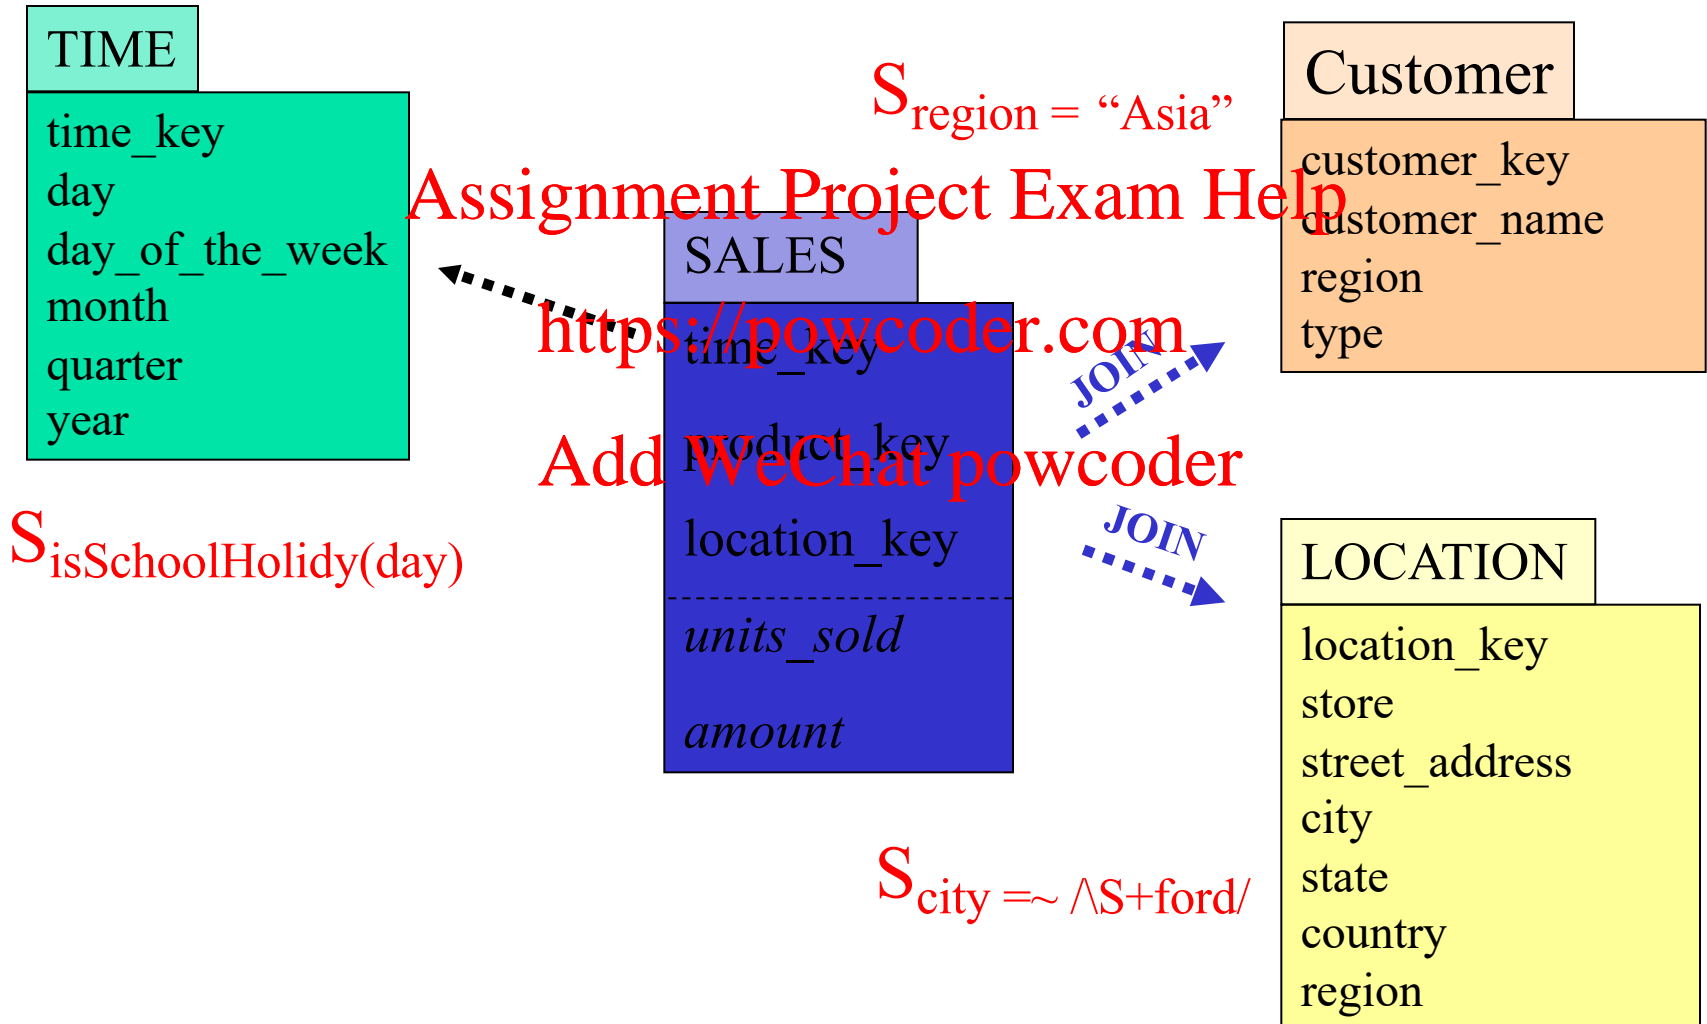
city = Sydney
city = Randwick
city = Kingsford
city = Coogee

Sales

<i>R102</i>	1
<i>R117</i>	1
<i>R118</i>	1
<i>R124</i>	1

Kingsford → R102,R117,R118,R124

Q3: Arbitrary selections on Dimensions




Star Query and Star Join (Cont.)

Usually only part of the dim tables because of the selection predicates

Time Customer Loc

time		Cust		loc
1	X	10	X	100
2		20		200



Sales

time	cust	loc	sold
1	10	100	7
1	10	150	13
1	20	150	2
2	20	200	16
...
1000	2000	500	86

*millions
of tuples*

<https://powcoder.com>

Add WeChat powcoder

time	cust	loc
1	10	100
1	10	200
1	20	100
1	20	200
2	10	100
2	10	200
2	20	100
2	20	200

*thousands
of tuples*

$\text{Sales} \triangleright \triangleleft \sigma_1(\text{Time}) \triangleright \triangleleft \sigma_2(\text{Cust}) \triangleright \triangleleft \sigma_3(\text{Loc}) \rightarrow$
 $\text{Sales} \triangleright \triangleleft (\sigma_1(\text{Time}) \times \sigma_2(\text{Cust}) \times \sigma_3(\text{Loc}))$

Q4: Coarse-grain Aggregations

- “Find total sales per customer type in our stores in Europe”
- Join-index will prune $\frac{3}{4}$ of the data (uniform sales), but the remaining $\frac{1}{4}$ is still large (several millions transactions)

- Index is unclustered

- High-level aggregations are expensive!!!!

⇒ Long Query Response Times

⇒ Pre-computation is necessary

⇒ Pre-computation is most beneficial

LOCATON

region

country

state


city

store

Cuboids = GROUP BYs

- Multidimensional aggregation = selection on corresponding cuboid

$GB_{(type, city)}(\sigma_1(Year) \bowtie \sigma_2(Brand) \bowtie \sigma_3(Loc))$



σ_1 selects some Years, σ_2 selects some Brands,
 σ_3 selects some Cities

$GB_{(type, city)}(\sigma_{1,2,3}(Cuboid(Year, Type, City)))$

- Materialize some/all of the cuboids
 - A complex decision involving cuboid sizes, query workload, and physical organization

Two Issues

- How to store the materialized cuboids?
- How to compute the cuboids efficiently?

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

CUBE BY in ROLAP

Sales		Product				
		1	2	3	4	ALL
Store	1	454	-	-	925	1379
	2	468	800	-	-	1268
	3	296	-	536	-	832
	4	652	-	540	745	1937
	ALL	1870	800	780	1170	5120

4 Group-bys here:
 (store,product)
 (store)
 (product)
 ()

- Need to write 4 queries!!!
- Compute them independently

Store	Product_key	sum(amtout)
1	1	454
1	4	925
2	1	468
2	2	800
3	1	296
3	3	240
4	1	625
4	3	240
4	4	745
1	ALL	1379
2	ALL	1268
3	ALL	536
4	ALL	1937
ALL	1	1870
ALL	2	800
ALL	3	780
ALL	4	1670
ALL	ALL	5120

SELECT LOCATION.store, SALES.product_key, SUM (amount)

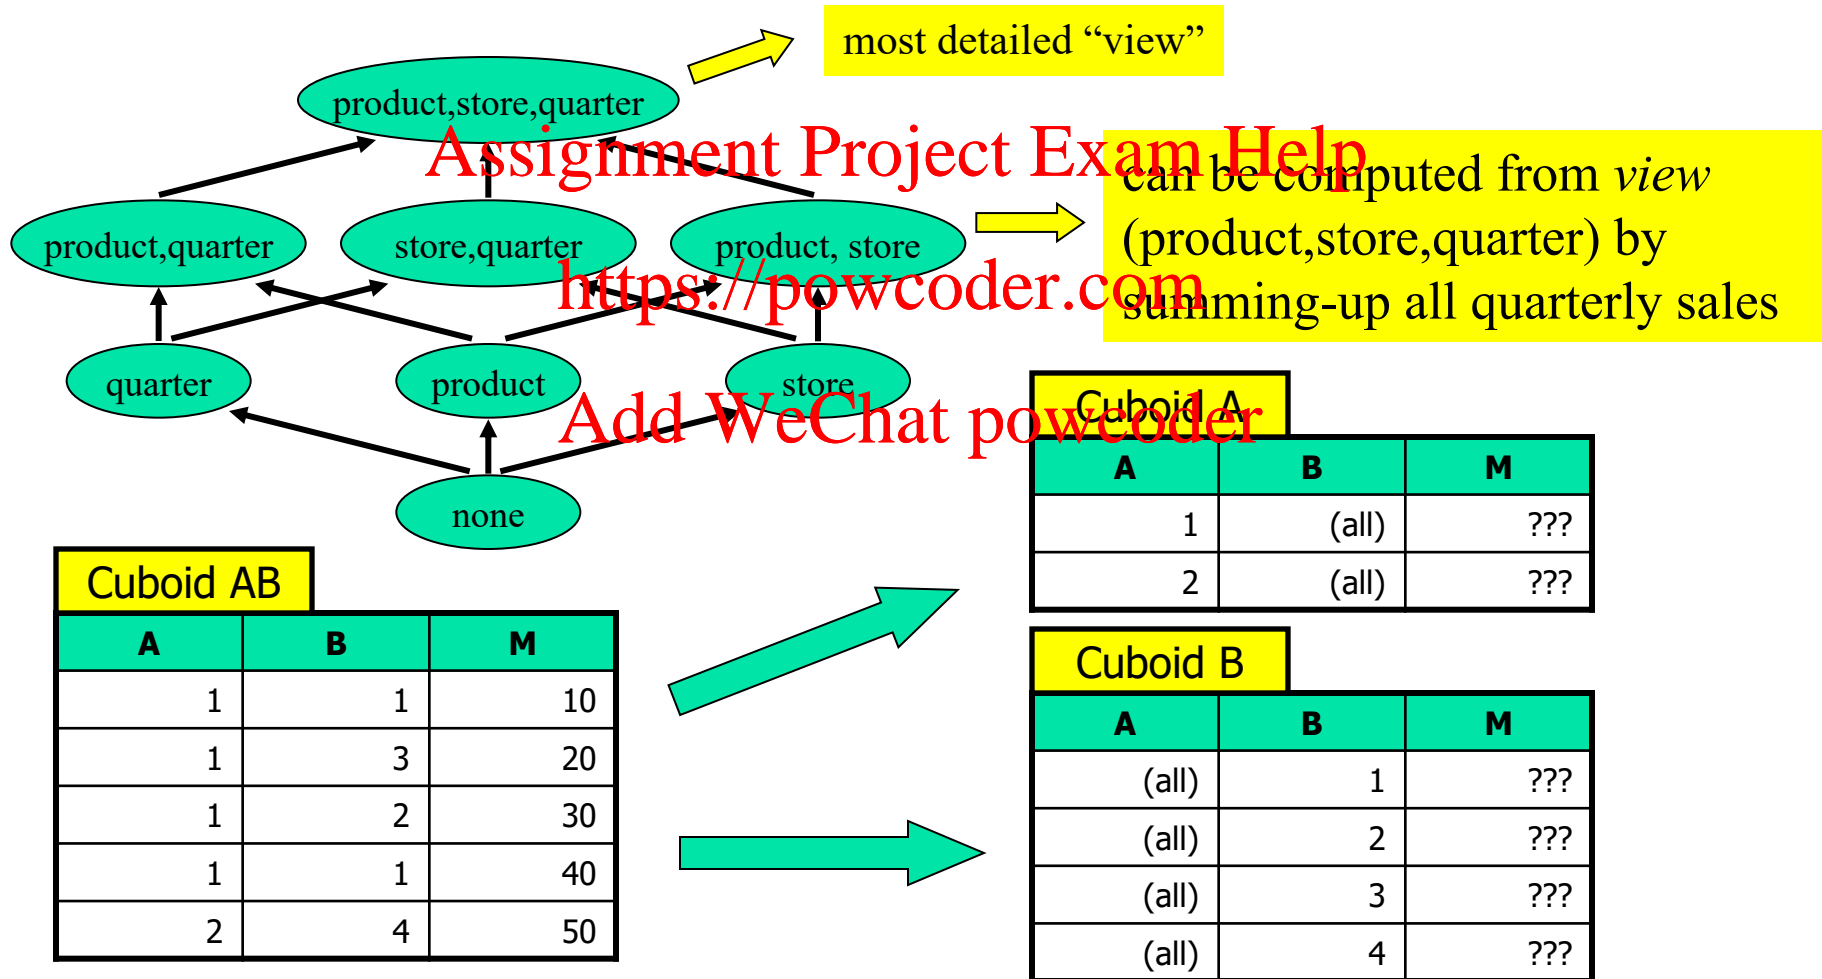
FROM SALES, LOCATION

WHERE SALES.location_key=LOCATION.location_key

CUBE BY SALES.product_key, LOCATION.store

Top-down Approach

- Model dependencies among the aggregates:



Bottom-Up Approach (BUC)

- BUC (Beyer & Ramakrishnan, SIGMOD'99)


- Ideas

- Compute the cube from bottom up
 - Divide-and-conquer


- A simpler recursive version:

- BUC-SR

A	B	...
1	1	...
1	3	...
1	2	...
1	1	...
2



...



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Understanding Recursion /1

- Powerful computing/problem-solving techniques

Assignment Project Exam Help

- Examples

- Factorial:

<https://powcoder.com>

- $f(n) = 1$, if $n = 1$
- $f(n) = f(n-1) * n$, if $n \geq 1$

Add WeChat powcoder

- Quick sort:

- $\text{Sort}([x]) = [x]$
- $\text{Sort}([x_1, \dots, \text{pivot}, \dots, x_n]) = \text{sort}[\text{ys}] ++ \text{sort}[\text{zs}]$, where

$\text{ys} = [x \mid x \text{ in } x_i, x \leq \text{pivot}]$

$\text{zs} = [x \mid x \leftarrow x_i, x > \text{pivot}]$

$f(0) = 0! =$
???

List comprehension
in Haskell or
python

Understanding Recursion /2


- Let $C(n, m)$ be the number of ways to select m balls from n numbered balls

- Show that $C(n, m) = C(n-1, m-1) + C(n-1, m)$

entire solution space

- Example: $m = 3, n = 5$

- Consider any ball in the 5 balls, e.g., 'd'



?i in { a b c d e }

Key Points

- Sub-problems need to be “**smaller**”, so that a simple/trivial boundary case can be reached
- Divide-and-conquer
 - There may be multiple ways the entire solution space can be divided into **disjoint** sub-spaces, each of which can be conquered **recursively**.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Geometric Intuition /1

- Reduce Cube(in 2D) to Cube(in 1D)

Assignment Project Exam Help

	b1	b2	b3	
a1	M11	M12	M13	[Step 1]
a2	M21	M22	M23	[Step 1]
	[Step 2]	[Step 2]	[Step 2]	[Step 3]

<https://powcoder.com>

Add WeChat powcoder

	b1	b2	b3	*
[a1] ×	M11	M12	M13	[Step 1]
[a2] ×	M21	M22	M23	[Step 1]
[*] ×	[Step 2]	[Step 2]	[Step 2]	[Step 3]


Geometric Intuition /2

- Reduce Cube(in 3D) to Cube(in 2D)


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Geometric Intuition /3






- Reduce Cube(in 3D) to Cube(in 2D)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Algebraic Derivation

- How to compute n -dim cube on $(n+1)$ -dim base cuboid (array)?
 - What do output tuples look like?
- How to compute $(n+1)$ -dim cube on $(n+1)$ -dim base cuboid (array)?
 - What else do we need?

[{r1-r5}, ABC]

[{r1-r5}, BC]

	A	B	C	M
r1	1	1	1	10
r2	1	1	2	20
r3	1	2	1	30
r4	1	3	1	40
r5	2	1	1	50

BUC-SR (Simple Recursion)*

- BUC-SR(data, dims)
 - If (dims is empty)
 - Output (sum(data))
 - Else
 - Dims = [dim1, rest_of_dims]
 - For each distinct value *v* of dim1
 - slice_*v* = slice of data on “dim1 = *v*”
 - BUC-SR(slice_*v*, rest_of_dims)
 - data' = Project(data, rest_of_dims)
 - BUC-SR(data', rest_of_dims)

Boundary case:
data is essentially a list
of measure values

General case:
1) Slice on dim1. Call
BUC-SR recursively for
each slice

2) Project out dim1, and
call BUC-SR on it
recursively

Example

	1	2	3	*
1	30	30	40	100
2	50			50
*	80	30	40	150

[{r1-r4}, B]

B	M
1	10
1	20
2	30
3	40

Internal Output

B	M
1	30
2	30
3	40
*	100

Output

A	B	M
1	1	30
1	2	30
1	3	40
1	*	100

Assignment Project Exam Help

<https://www.powcoder.com>

Add WeChat powcoder

[{r1-r5}, AB]

A	B	M
1	1	10
1	1	20
1	2	30
1	3	40
2	1	50

[{r5}, B]

B	M
1	50

B	M
1	50
*	50

A	B	M
2	1	50
2	*	50

[{r1'-r5'}, B]

B	M
1	80
2	30
3	40

B	M
1	80
2	30
3	40
*	150

A	B	M
*	1	80
*	2	30
*	3	40
*	*	150

r1
r2
r3
r4
r5

Try a 3D-Cube by Yourself

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

[{r1-r5}, ABC]

	A	B	C	M
r1	1	1	1	10
r2	1	1	2	20
r3	1	2	1	30
r4	1	3	1	40
r5	2	1	1	50

MOLAP

- (Sparse) array-based multidimensional storage engine
- Pros:
 - small size (esp. for dense cubes)
 - fast in indexing and query processing
- Cons:
 - scalability
 - conversion from relational data

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Multidimensional Array

$$f(\text{time}, \text{item}) = 4 * \text{time} + \text{item}$$

time	item	dollars_sold
Q1	home entertainment	605
Q2	home entertainment	680
Q3	home entertainment	812
Q4	home entertainment	927
Q1	computer	825
Q2	computer	952
Q3	computer	1023
Q4	computer	1038
Q1	phone	14
Q2	phone	31
Q3	phone	30
Q4	phone	38
Q1	security	400
Q2	security	512
Q3	security	501
Q4	security	580



Step 1

Mappings

time	value
Q1	0
Q2	1
Q3	2
Q4	3

item	value
home entertainment	0
computer	1
phone	2
security	3

time	item	dollars_sold
0	0	605
1	0	680
2	0	812
3	0	927
0	1	825
1	1	952
2	1	1023
3	1	1038
0	2	14
1	2	31
2	2	30
3	2	38
0	3	400
1	3	512
2	3	501
3	3	580

offset

0
4
8
12
1
5
9
13
2
6
10
14
3
7
11
15

Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder

Multidimensional Array

Step 3: If **dense**, only need to store sorted slots

offset	dollars_sold
0	605
1	825
2	14
3	105
4	680
5	952
6	31
7	512
8	812
9	1023
10	30
11	501
12	927
13	1038
14	38
15	580



- Think: how to decode a slot?

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Dense MD array
605
825
14
400
680
952
31
512
812
1023
30
501
927
1038
38
580

$$f(\text{time}, \text{item}) = 4 * \text{time} + \text{item}$$

The Sparse Case

time	item	dollars_sold
Q1	home entertainment	605
/* same table but with many rows deleted to make it sparse */		
Q4	security	580

Step 1

Mappings	
time	value
Q1	0
Q2	1
Q3	2
Q4	3

item	value
home entertainment	0
computer	1
phone	2
security	3

time	item	dollars_sold	offset
0	0	605	0
/* same table but with many rows deleted to make it sparse */			
3	3	580	15

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Multidimensional Array

Choice 1

offset	dollars_sold
0	605
15	580

Choice 2

Dense MD array
605
-
-
-
-
-
-
-
-
-
-
-
-
-
-
580

- Think: how to decode a slot?
- Multidimensional array is typically sparse
- Use sparse array (i.e., offset + value)
- Could use chunk to further reduce the space
- Space usage:
 - $(d+1)*n*4$ vs $2*n*4$
- HOLAP:
 - Store all non-base cuboid in MD array
 - Assign a value for ALL

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder